

FlexOS: Making OS Isolation Flexible

H. Lefeuvre, V-A. Badoiu, S.L. Teodorescu, P. Olivier,
T. Mosnoi, R. Deaconescu, F. Huici, C. Raiciu


HotOS 2021, 1 June–3 June 2021



Current OS Designs

OS security/isolation strategies are **fixed** at design time!

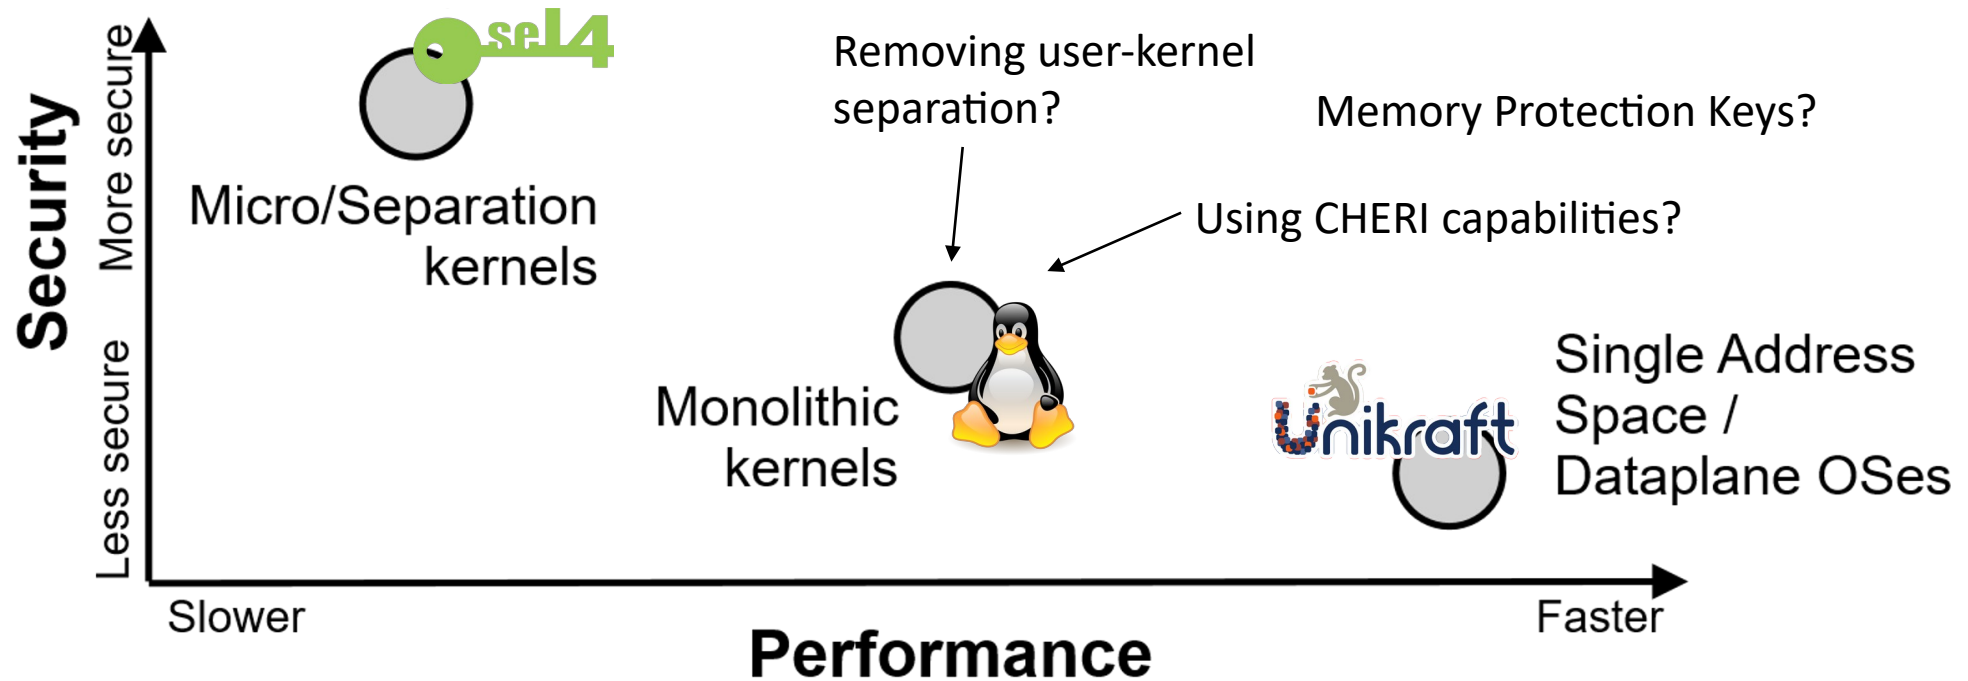
Isolation granularity, underlying mechanisms, data sharing strategies (copy/share)



Current OS Designs


OS security/isolation strategies are **fixed** at design time!

Isolation granularity, underlying mechanisms, data sharing strategies (copy/share)



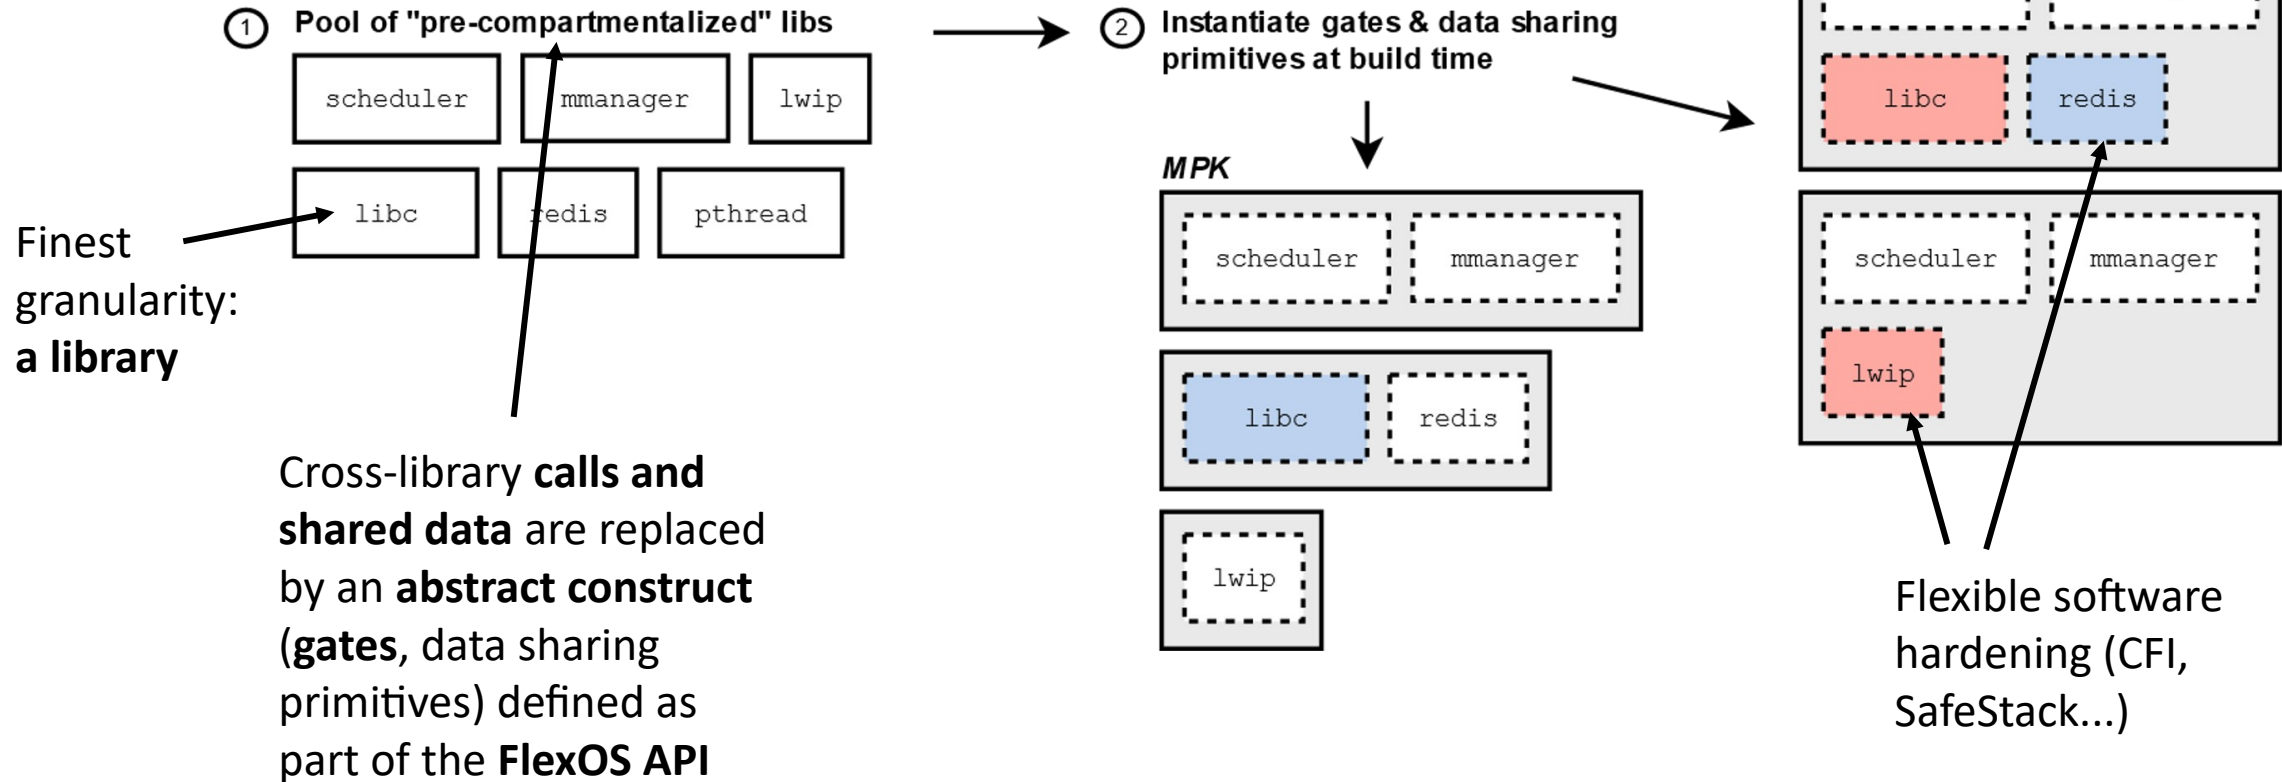
FlexOS: Flexible Isolation

Decouple security/isolation decisions from the OS design

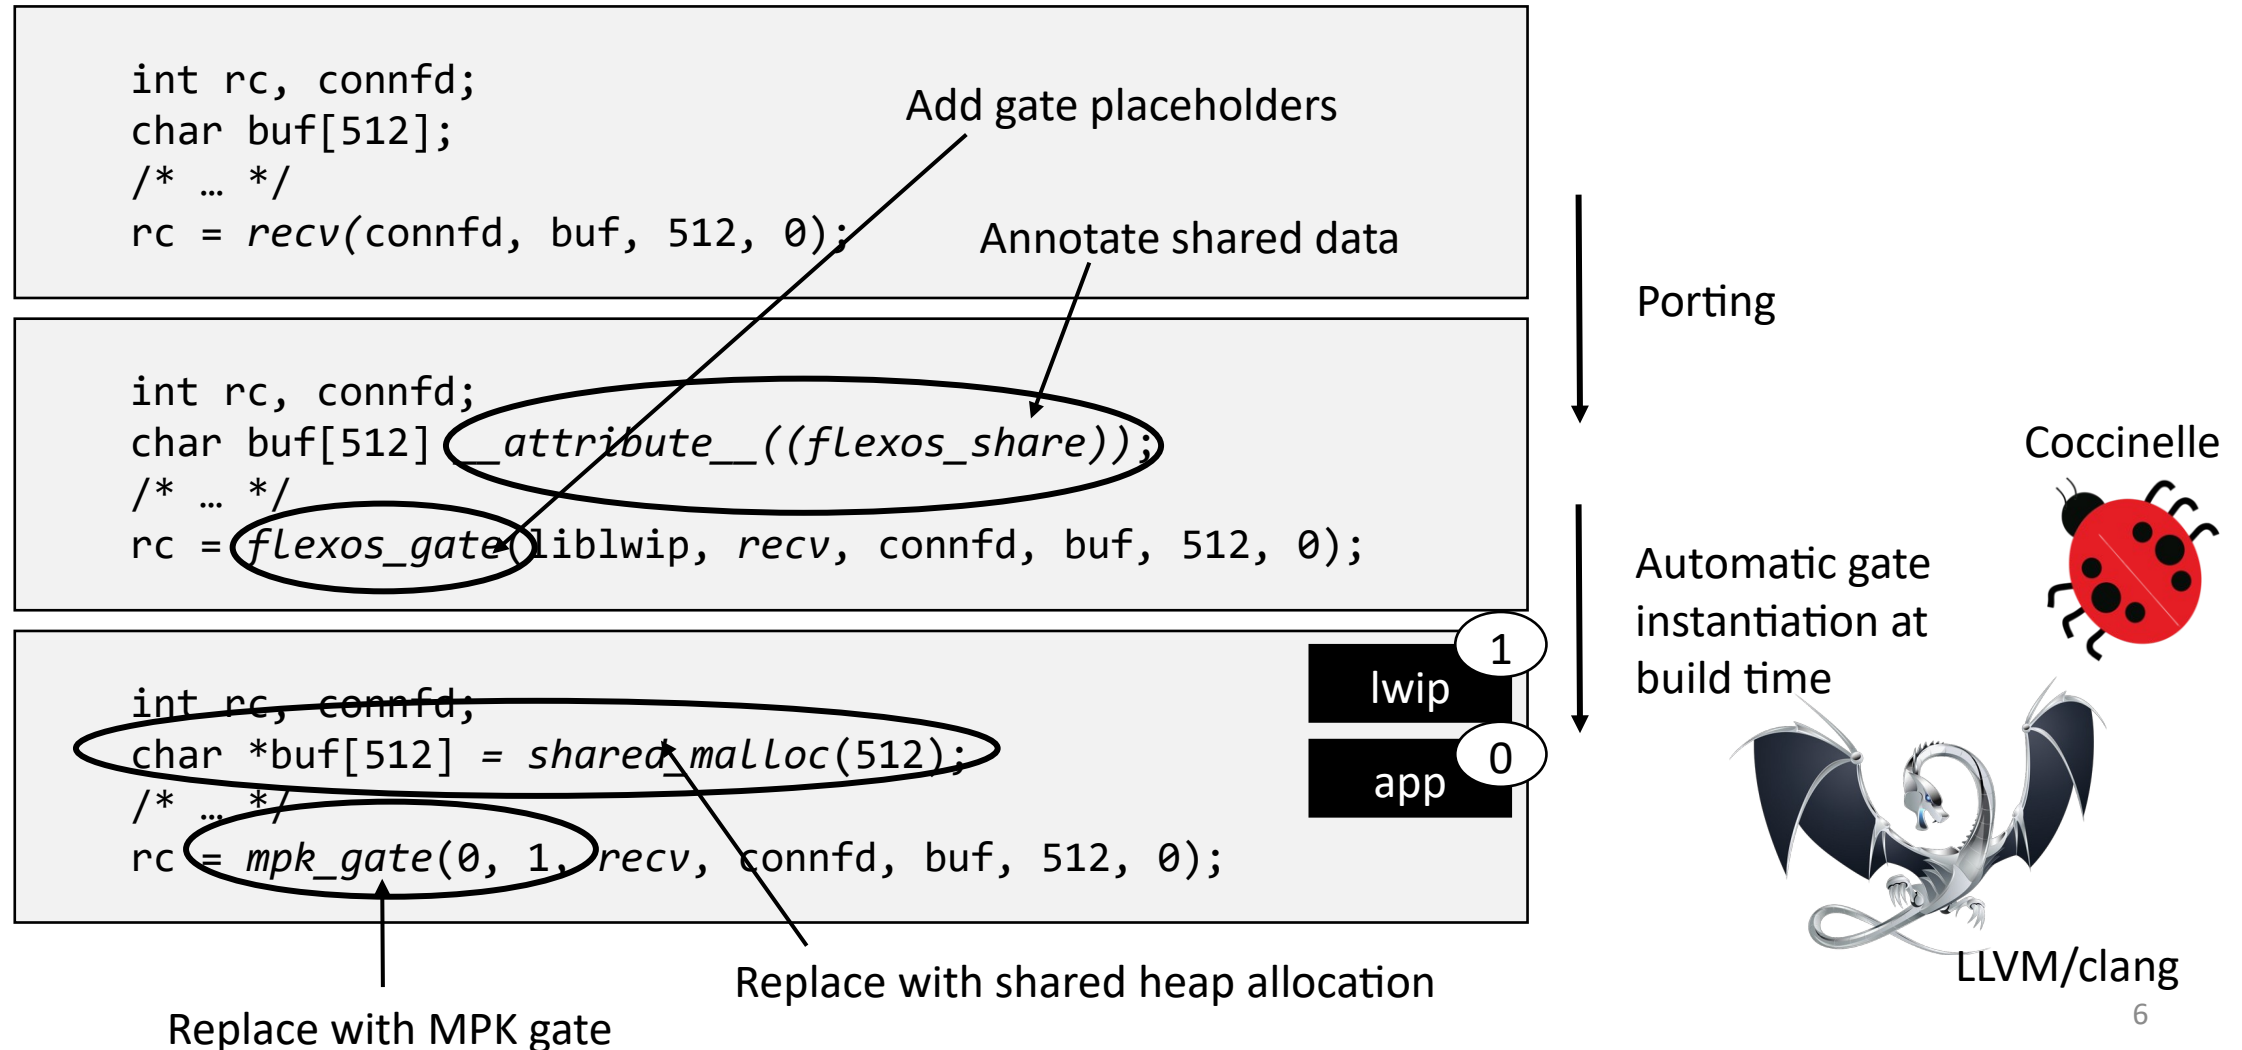


FlexOS 101: Flexible Builds

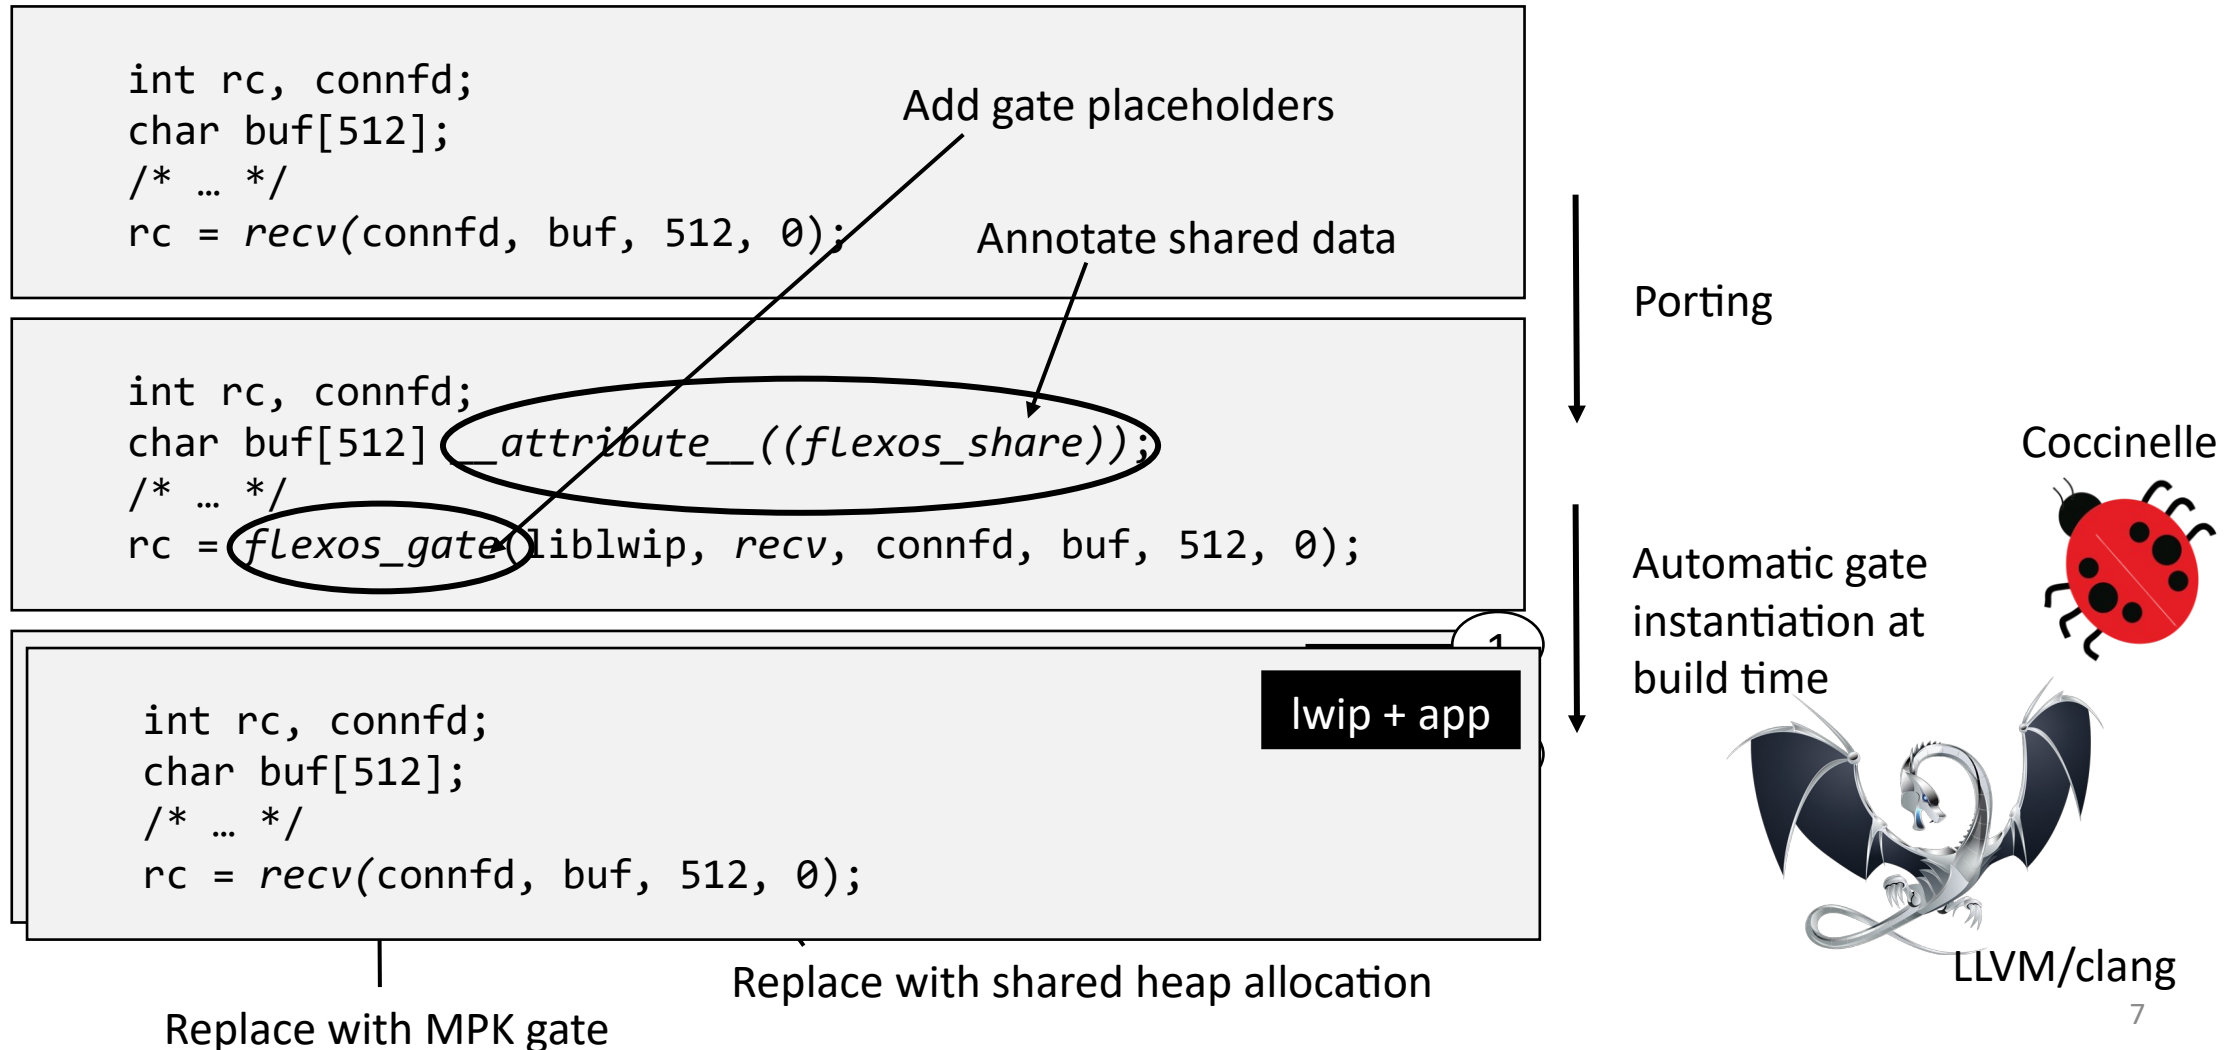
Based on a **highly modular LibOS design** (Unikraft)



FlexOS 101: Compartmentalization API



FlexOS 101: Compartmentalization API



FlexOS 101: Design Space

Huge design space! Not all configurations make sense from...

- a **security** perspective
- a **performance** perspective

Security perspective: how can we **guarantee that the properties of components hold?**

Enrich the OS with a tool that

1. selects configurations so that **properties hold**
2. **further prunes based on security and performance**

FlexOS 101: Do my properties hold?

Each component has a set of properties and expectations from other components!

Describe them with a DSL

Formally verified scheduler

```
[Memory access] Read(Own,Shared); Write(Own,Shared)
[Call] alloc::malloc, alloc::free
[API] thread_add (...); thread_rm(...); yield(...)
[Requires] *(Read,Own), *(Write,Shared),
            *(Call, thread_add), *...
```

What is the component's API?


Requirements on other components?

What memory can this component access?

What functions can this component call?

Plain-C NW stack

```
[Memory access] Read(*); Write(*)
[Call] *
```



FlexOS 101: Do my properties hold?

Using the DSL, we can determine **pair-wise compatibility** such that **properties hold**

Plain-C NW stack

```
[Memory access] Read(*); Write(*)  
[Call] *
```


Formally verified scheduler

```
[Memory access] Read(Own,Shared); Write(Own,Shared);  
[Call] alloc::malloc, alloc::free  
[API] thread_add (...); thread_rm(...); yield(...)  
[Requires] *(Read,Own), *(Write,Shared),  
*(Call, thread_add), *...
```



FlexOS 101: Do my properties hold?

Using the DSL, we can determine **pair-wise compatibility** such that **properties hold**



Prototype

Implementation **on top of Unikraft**

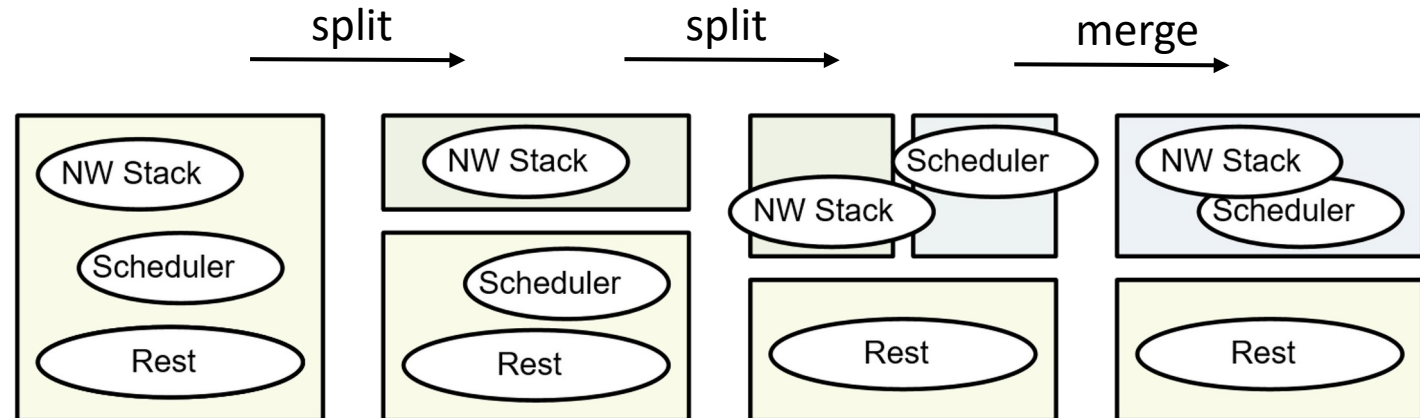
Gate implementations for **Intel MPK** and **VM/EPT**

Port of the **network stack** (lwip) and **scheduler** (uksched)

Prototype

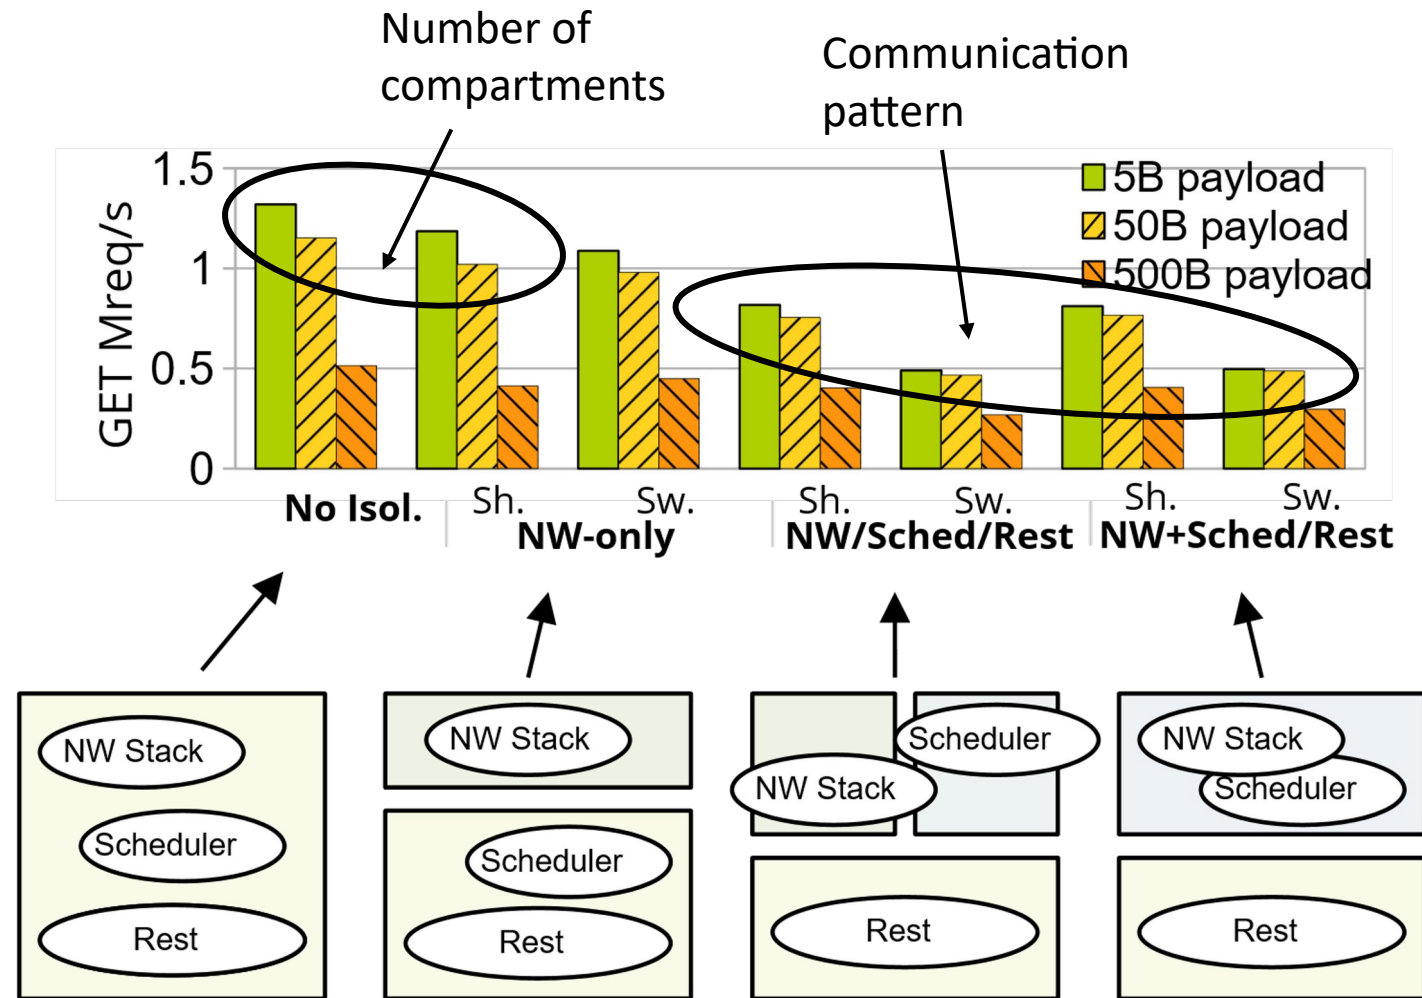
Here: **MPK** and **Redis**

4 configurations, up to 3 compartments



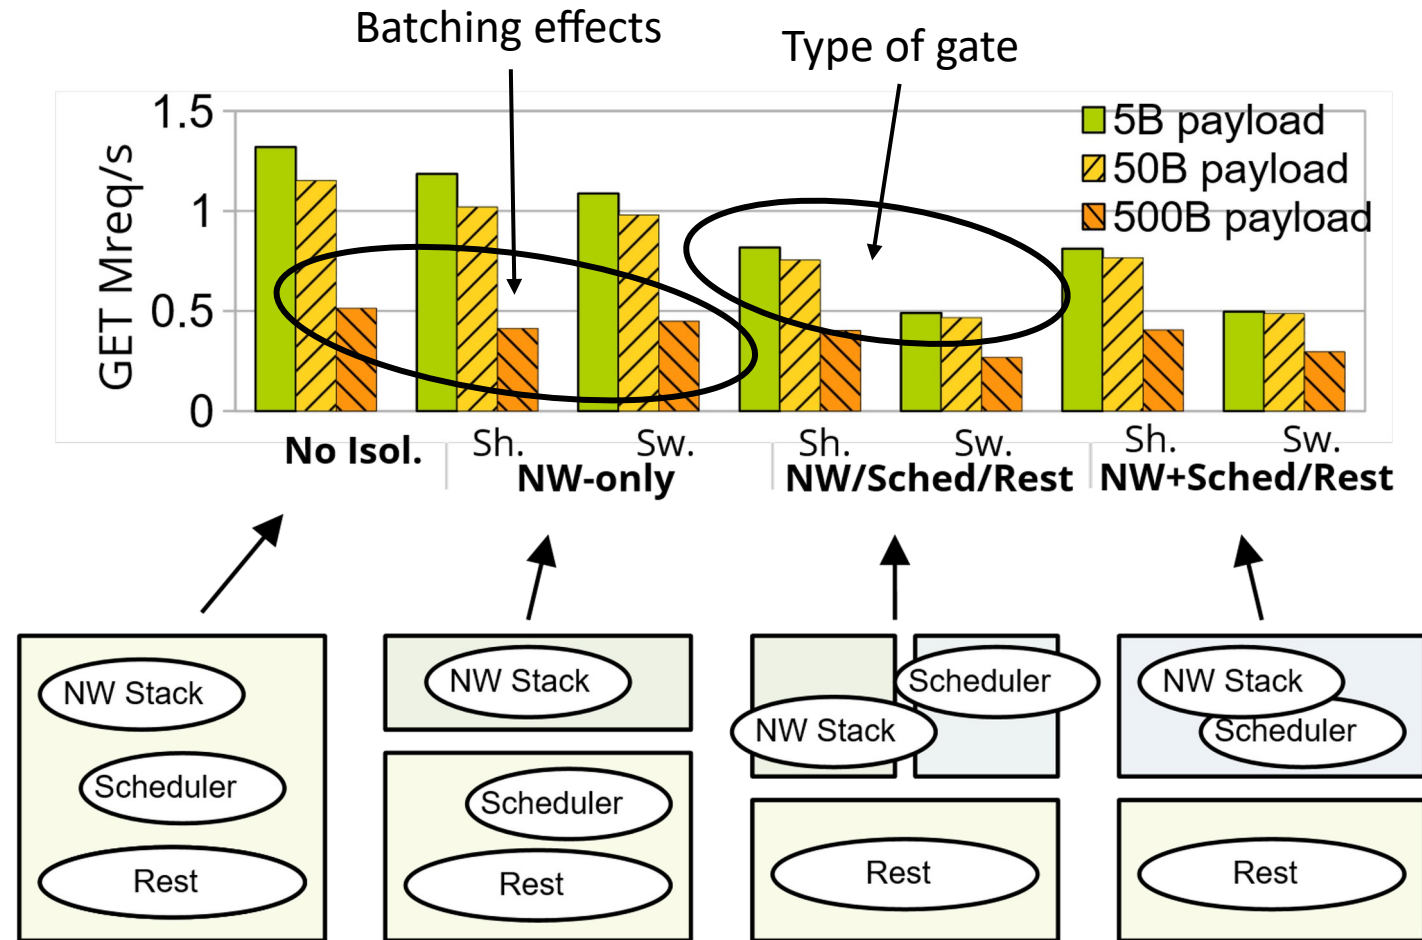
Prototype

Here: **MPK and Redis**



Prototype

Here: **MPK and Redis**

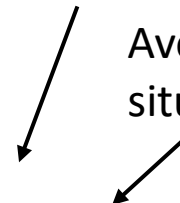


Some open questions...

1. **Security is not only a matter of isolation**

- When APIs are **designed as a trust boundary**, swapping the isolation mechanism is easy!
- *What if the API wasn't developed as a trust boundary?*

Carefully check arguments
Avoid confused deputy situations

Two arrows originate from the text. One arrow points from 'Carefully check arguments' to the first bullet point. The other arrow points from 'Avoid confused deputy situations' to the same bullet point.

2. **How to minimize porting effort?**

- FlexOS requires porting of app/kernel libraries
- *How to automate the process?*
- *How to make sure that DSL metadata is correct?*

In a Nutshell

There is a **need for isolation flexibility**

- OS Specialization, hardware heterogeneity
- or even breaking primitives!

Current approaches: **one isolation approach at design time**

Decouple isolation from the OS design:

- Make isolation decisions at **build time**
- Automatically **explore performance v.s. security trade-offs**

